

Deep Learning – The Big Picture

C. Heumann, M. Aßenmacher, D. Schalk

April 2, 2020



Brief History of Deep Learning

Deep Learning Timeline - 1

1943

Warren S. McCulloch and Walter Pitts tried to understand how the brain could produce highly complex patterns by using many basic cells that are connected together.

1943

Deep Learning Timeline - 1

1943

Warren S. McCulloch and Walter Pitts tried to understand how the brain could produce highly complex patterns by using many basic cells that are connected together.

1943

1950

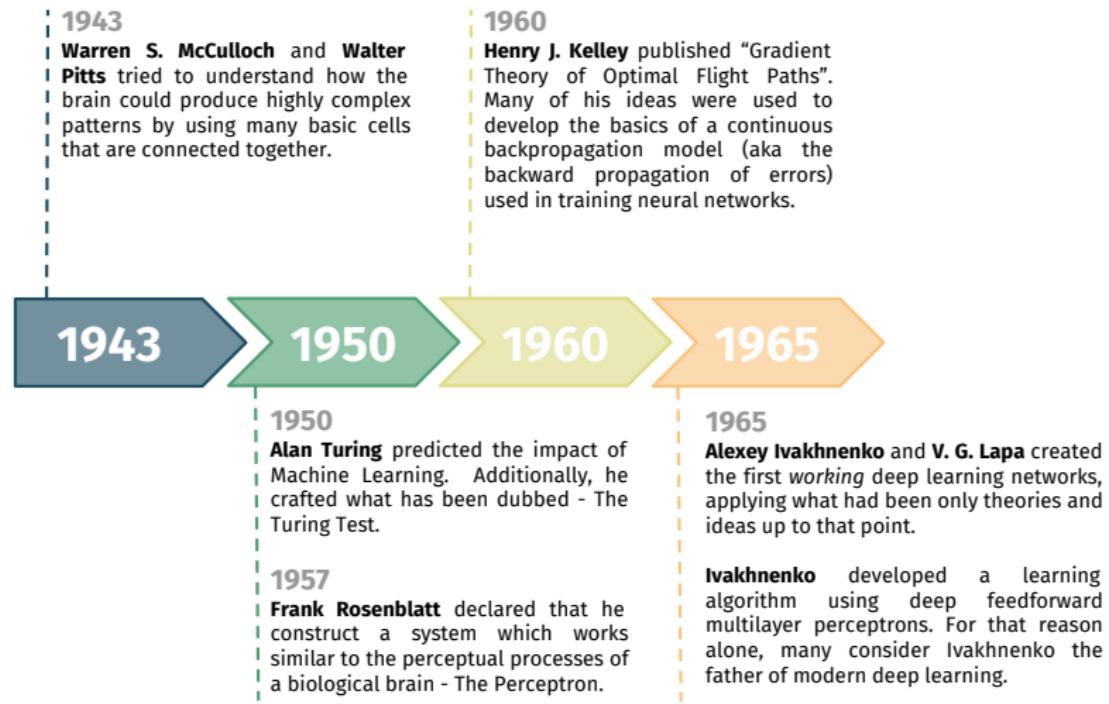
1950

Alan Turing predicted the impact of Machine Learning. Additionally, he crafted what has been dubbed - The Turing Test.

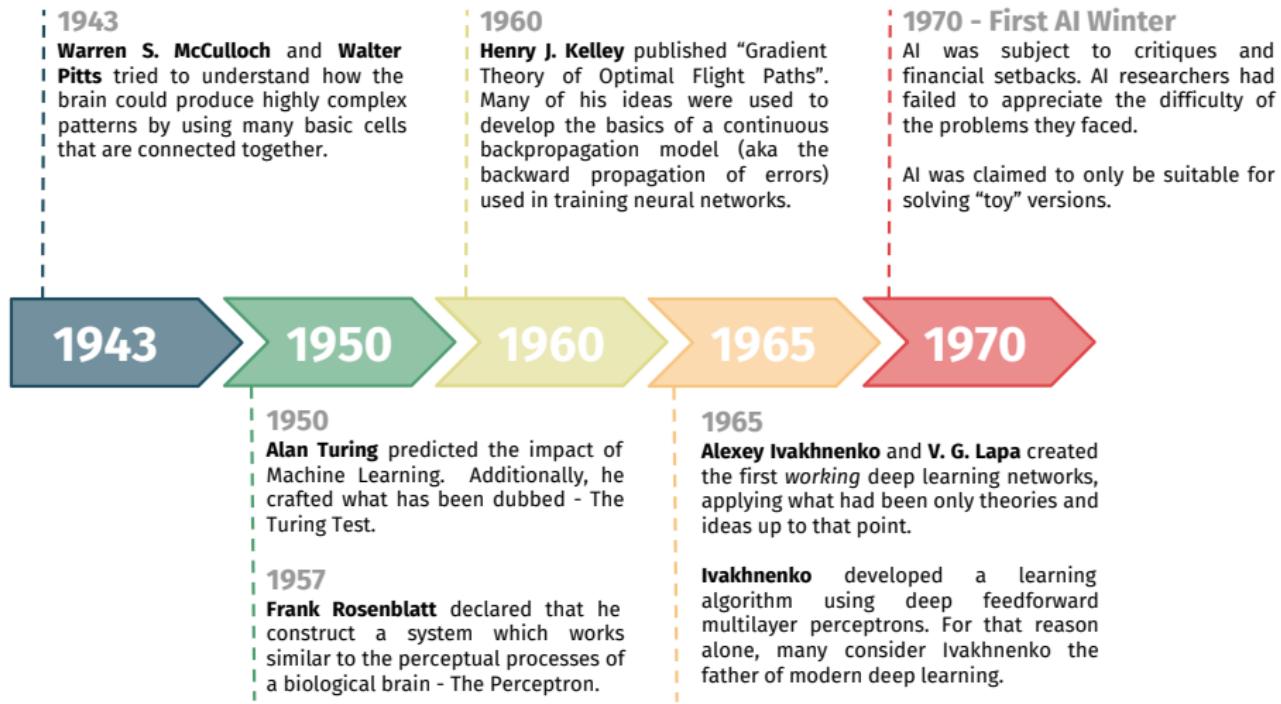
1957

Frank Rosenblatt declared that he construct a system which works similar to the perceptual processes of a biological brain - The Perceptron.

Deep Learning Timeline - 1



Deep Learning Timeline - 1



Deep Learning Timeline - 2

1980

Kunihiko Fukushima creates Neocognitron, an artificial neural network that learned how to recognize visual patterns.

His work led to the development of the first convolutional neural networks, which are based on the visual cortex organization found in animals.

1980

Deep Learning Timeline - 2

1980

Kunihiko Fukushima creates Neocognitron, an artificial neural network that learned how to recognize visual patterns.

His work led to the development of the first convolutional neural networks, which are based on the visual cortex organization found in animals.

1980

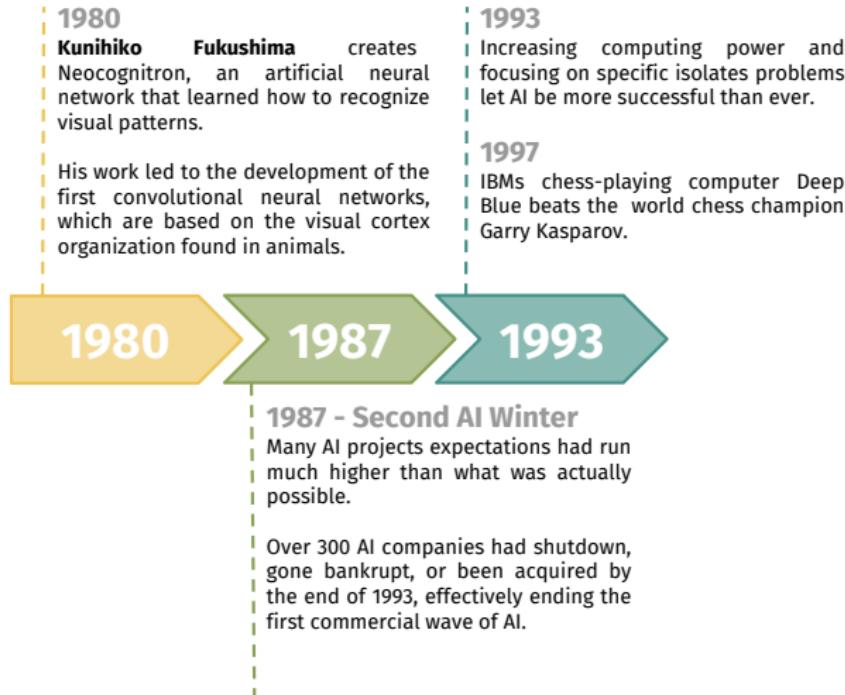
1987

1987 - Second AI Winter

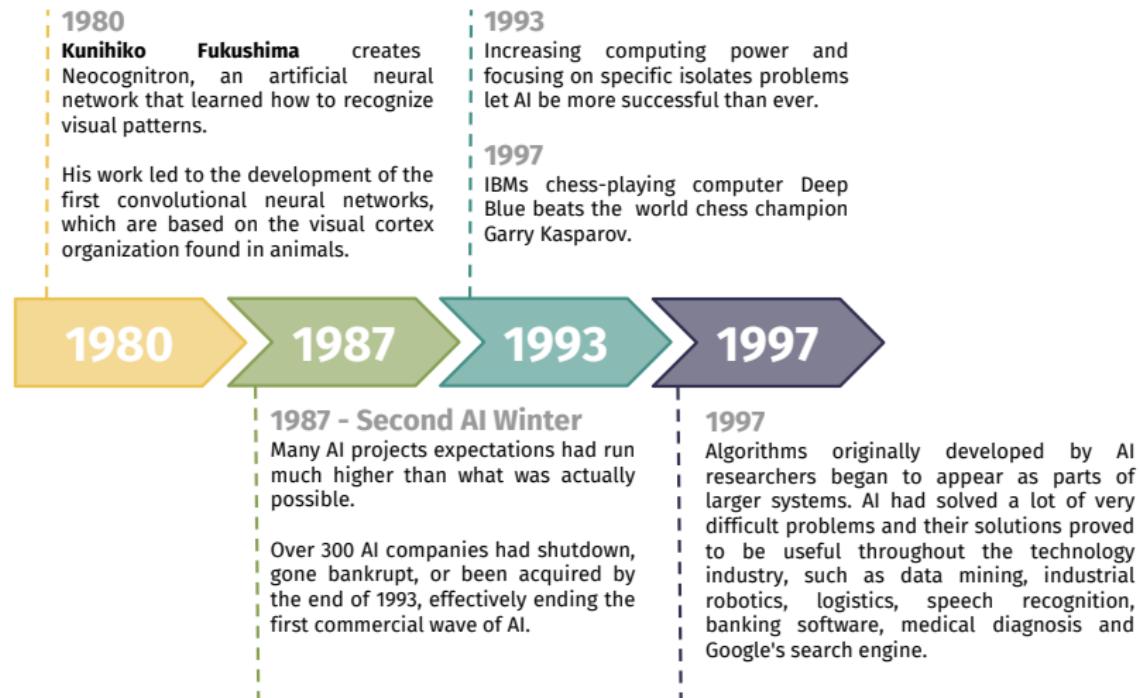
Many AI projects expectations had run much higher than what was actually possible.

Over 300 AI companies had shutdown, gone bankrupt, or been acquired by the end of 1993, effectively ending the first commercial wave of AI.

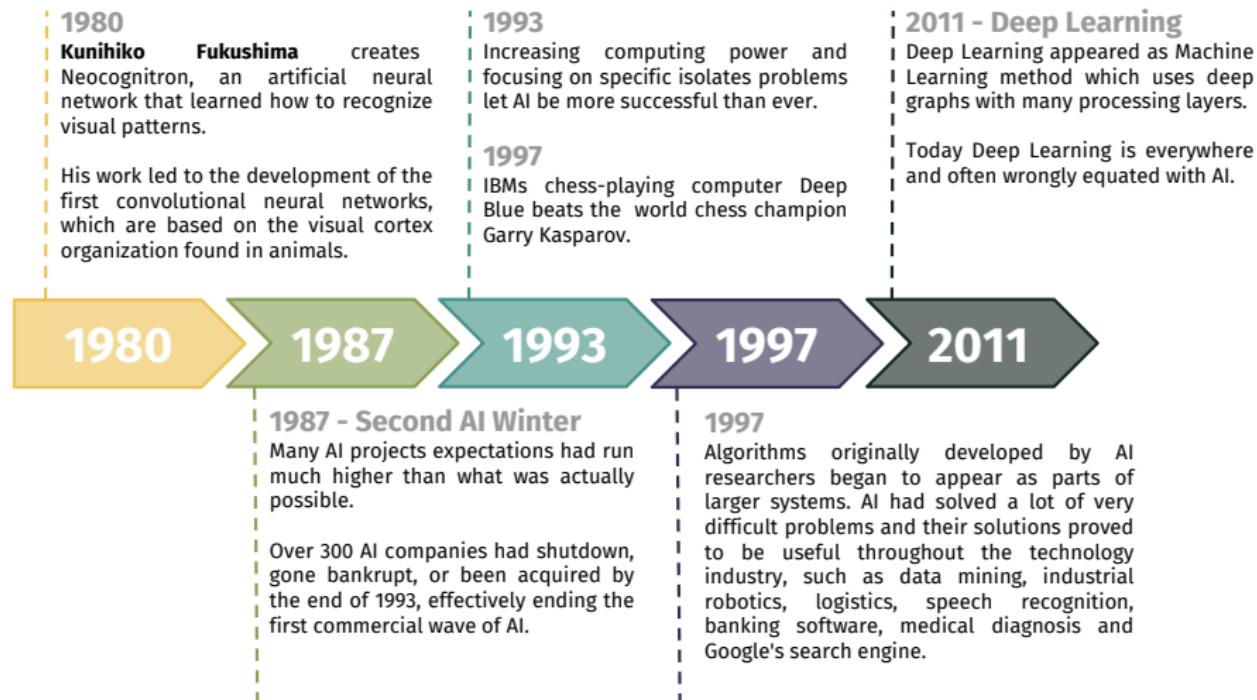
Deep Learning Timeline - 2



Deep Learning Timeline - 2



Deep Learning Timeline - 2



Deep Learning Timeline - 3

2013 - word2vec

Tomas Mikolov et al. publish four papers on vector representations of words constituting the *word2vec* framework

This received very much attention as it revolutionized the way words were encoded for deep learning models in the field of NLP.

2013

Deep Learning Timeline - 3

2013 - word2vec

Tomas Mikolov et al. publish four papers on vector representations of words constituting the *word2vec* framework.

This received very much attention as it revolutionized the way words were encoded for deep learning models in the field of NLP.

2013

2014

2014 - Attention!

The concept of Attention is introduced by Bahdanau et al. as a complement to state-of-the-art neural network architectures.

It is (at first) just used in conjunction with RNNs for machine translation, but will kick-start a completely new path of Deep learning research in NLP in the years to come.

Deep Learning Timeline - 3

2013 - word2vec

Tomas Mikolov et al. publish four papers on vector representations of words constituting the *word2vec* framework

This received very much attention as it revolutionized the way words were encoded for deep learning models in the field of NLP.

2013

2014

2016

2016 - TPUs

Google developed a processing unit exceptionally suited for deep learning applications.

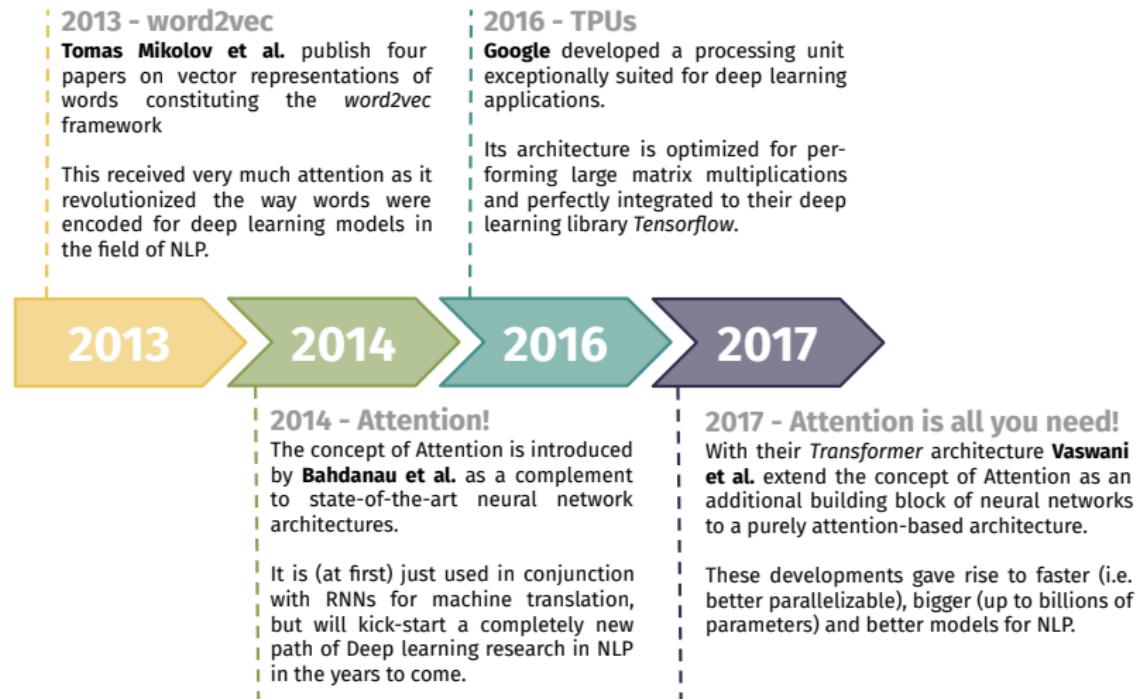
Its architecture is optimized for performing large matrix multiplications and perfectly integrated to their deep learning library *Tensorflow*.

2014 - Attention!

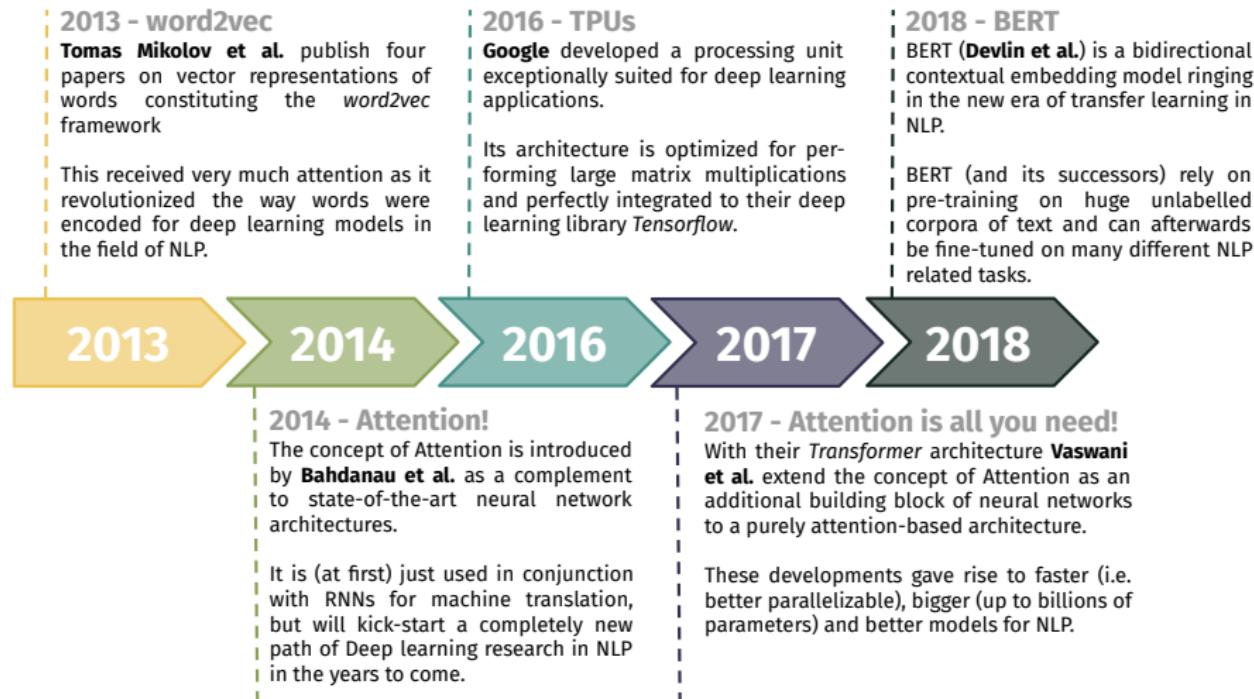
The concept of Attention is introduced by **Bahdanau et al.** as a complement to state-of-the-art neural network architectures.

It is (at first) just used in conjunction with RNNs for machine translation, but will kick-start a completely new path of Deep learning research in NLP in the years to come.

Deep Learning Timeline - 3



Deep Learning Timeline - 3



What is Possible With Deep Learning?

Deep Learning Applications

- Supervised learning tasks
- Object Recognition – Seeing:



- Speech Recognition and NLP – Hearing and understanding text:



Deep Learning Applications

- And know we try to teach them creativity:

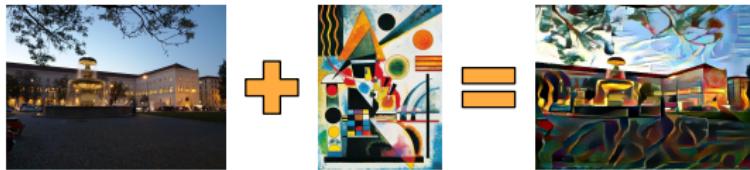
- **Music and Text Generation**

Example of generated text trained on Trump tweets:

"The Fake News Media is a great people of the president was a great people of the many people who would be a great people of the president was a big crowd of the statement of the media is a great people of the people of the statement of the people . . . "

... could be better.

- **Neural Style Transfer**

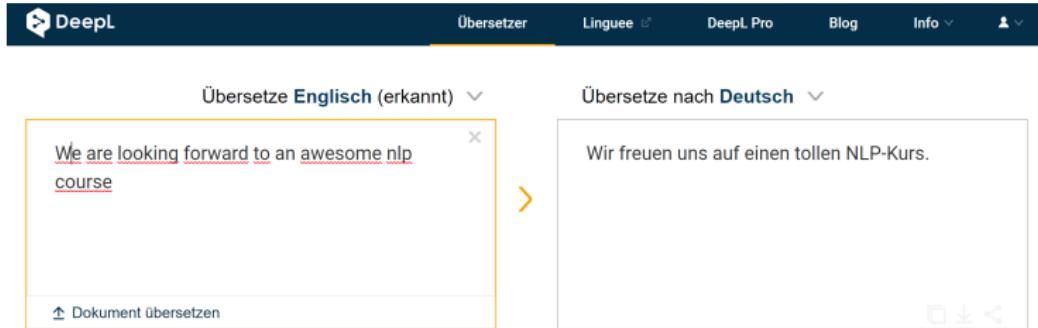


- ...

Some remarks on NLP

- The world's largest companies are investing heavily in NLP: Google, Facebook, Amazon, Apple, Microsoft, ...
- But still: a lot of (yet) unsolved problems and challenges
- Innovations are more and more open-sourced
- There's a variety of different applications of NLP in everyday life you might not be aware of
- Once you see them, it soon becomes obvious

Neural Machine Translation

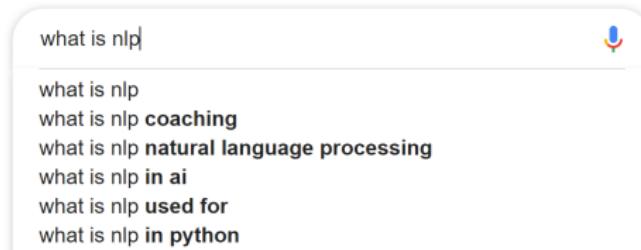


- Two different NLP tasks are hiding in this picture..
- **Underlying task:** Text generation

A brief History of Machine Translation

- **Rule-Based Machine Translation** (50s – 80s)
 - Dictionaries + Grammatical Rules
- **Example-Based Machine Translation** (80s – 90s)
 - First suggested by Makoto Nagao (1984)
 - Based on bilingual text corpora
- **Statistical Machine Translation** (90s – 10s)
 - Driven by IBM guys
- **Neural Machine Translation** (last few years)
 - Based on neural networks (LSTMs, Transformers)

Auto Completion



- You might also know this from your smartphone keyboard
- **Underlying task:** Predict the next word given the previous words (+ individual browsing history)

Spam filters



Dies ist ein Muss!

Das Angebot ist begrenzt, Details hier <https://hakingholess1975.blogspot.co.il/>

Viel Glück!

- Indeed, spam filters are often surprisingly simple systems
- **Underlying task:** Document classification (Naive Bayes or Logistic Regression)

Argument Mining

The screenshot shows the ArgumentText web application interface. At the top, there is a logo with a blue square icon containing a white triangle, followed by the text "ArgumentText". To the right of the logo are language selection ("EN") and search ("vegan diet") fields, and a magnifying glass search icon. Below the header is a navigation bar with tabs: "PRO/CON", "LIST", "WEIGHTS", and "DOCUMENTS". To the right of the navigation bar are "Filter" and "Sort By" buttons with dropdown menus. The main content area displays search results. A message at the top says "Found 87 arguments (55 pro; 32 con) in 16 documents (classified 437 sentences in 13.44 s)". Below this, there are two columns of arguments. The left column is labeled "PRO" and the right column is labeled "CON". Each row contains an argument statement, its source URL, and the date it was last updated. The first PRO argument is about a vegan diet providing physical benefits. The first CON argument is about the effects of eating raw fruit on health conditions like candida, chronic fatigue, diabetes, and cancer. The second PRO argument is about a naturally-occurring arthritis cure. The second CON argument is about vegans' environmental reasons for their diet.

| PRO | CON |
|--|--|
| A Vegan Diet Provides Physical Benefits In addition to disease prevention, a vegan diet provides additional physical health benefits to the body, which, helps a person, fight off cancer. 92.01% | This nutrition course also examines the effects of eating raw fruit on specific health conditions such as candida, chronic fatigue, diabetes, cancer, and the healthful management of weight. 94.81% |
| Cure Arthritis Naturally - Blue Heron Health News Vegan Diet Cure Rheumatoid Arthritis How to Vegan Diet Cure Rheumatoid Arthritis. 91.85% | Vegans for environmental reasons may not only lapse but may decide that an animal product has fewer adverse environmental consequences than non-animal products. 93.77% |

- Detection and ranking of arguments pro/contra a certain topic
- **Underlying task:** Document classification and/or Sequence tagging

Chatbots

Hallo! Schön, dass du hier bist. Was kann ich für dich tun?

Stelle mir deine Frage...



- We experience chatbots on a growing number of websites
- **Underlying task:** Natural language understanding / Text generation

Alexa, what is the capital of Australia?

- **Underlying task:** Question Answering / Information extraction

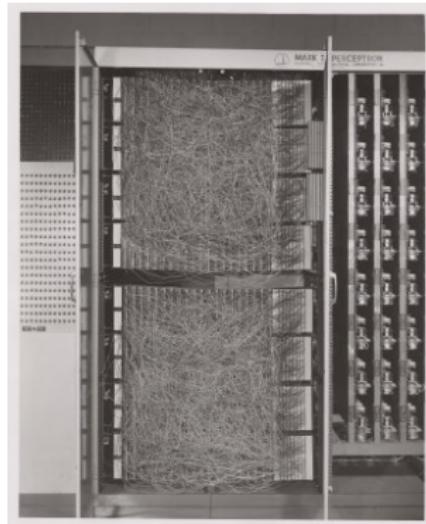
(Some) Further examples

- **Sentiment analysis / Text classification**
 - Determine, whether a document is positive/neutral/negative
- **Text Generation**
 - Example: *OpenAI GPT2*
- **Named Entity Recognition**
 - Search for certain entities (e.g. person, location, etc.)
- **Image Captioning**
 - Hybrid task (Computer Vision + NLP)
- **Dependency Parsing**
 - Determine the grammatical structure of text
- **Part-of-speech tagging**
 - Tagging words according to their role in a text

How Deep Learning Works

The Perceptron

- The perceptron was invented by Frank Rosenblatt 1957.

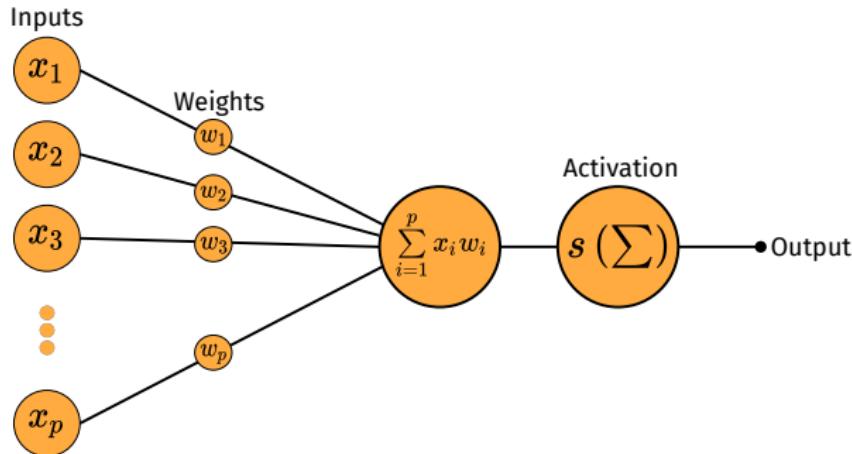


The Mark I Perceptron

- It is the basic computational unit for neural networks.

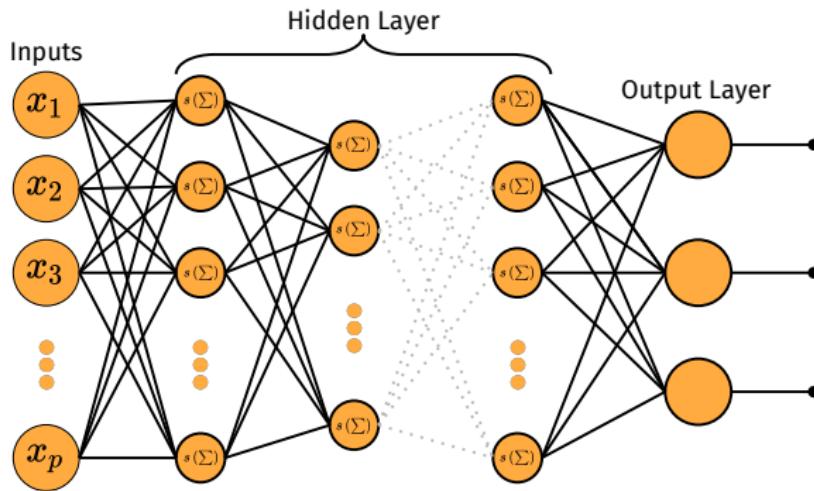
Singlelayer Perceptron

- Weighted sum of input values transformed by an activation function s
- If s is the sigmoid function $(1 + \exp \sum)^{-1}$, then the perceptron does exactly the same as the logistic regression



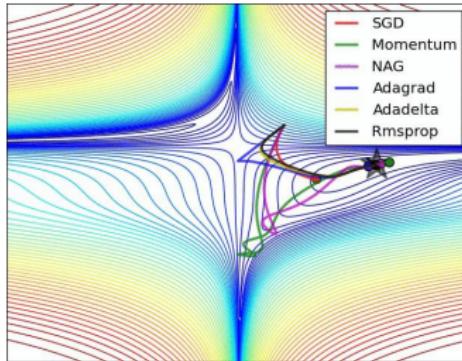
Multilayer Perceptron

- Stacking of multiple perceptrons
- Corresponds to stacking GLM models
- Number of parameter grows very fast
→ Optimization becomes more difficult



Optimizer

- Having that much parameter/weights to find, standard optimizer like Gradient Descent may fail
- Therefore, much effort was spent to get faster optimizer like adam, rmsprop, etc.:



Source: Ruder, S. (2016). An overview of gradient descent optimization algorithms. arXiv preprint arXiv:1609.04747.

Input Source Images

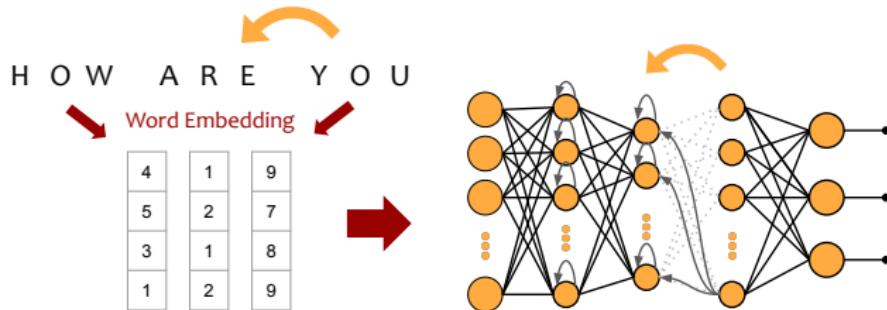
- Deep Learning is well known for handling more complex data structures like images, text, voice recordings etc.
 - When using images as input, one image becomes one observation:



- But, how to analyse images? We have to make use of the spatial structure, but how?

Input Source Text

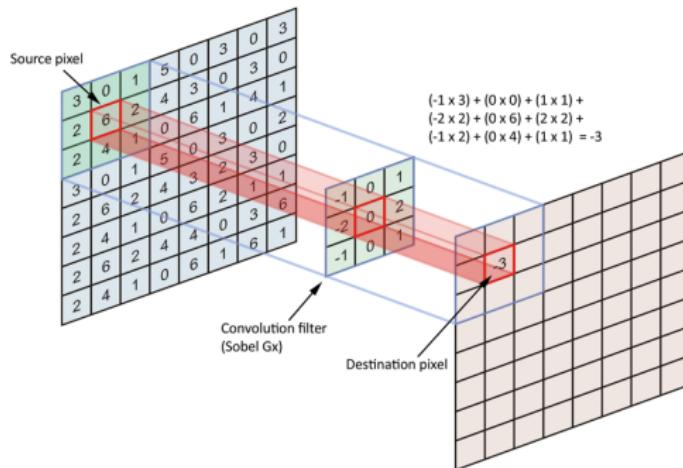
- Another typical data format are texts:



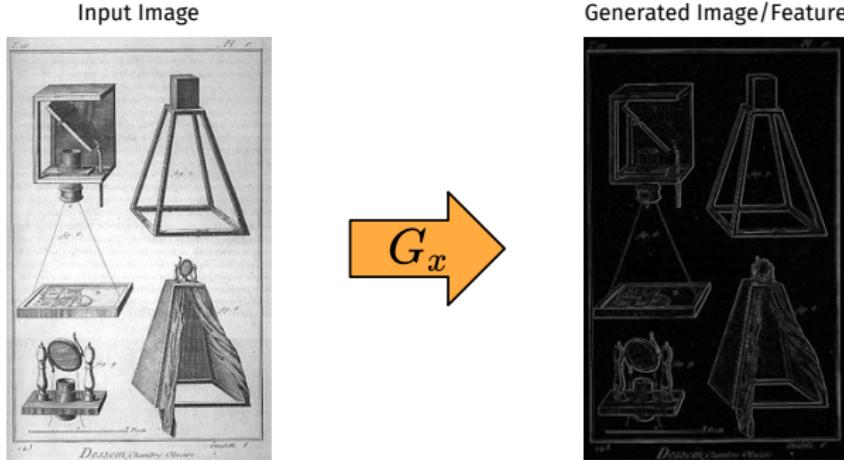
- But, how to analyse text sources? We also have to consider the recurrent dependencies of text, but how?

Convolution - As Operation

- Extraction of features of the input data:



Convolution - As Operation



Note: The recognition of edges and corners requires a multiple application of the operator.

Convolution - In Neural Networks

- Filters, like the Sobel filter, are obtained by thinking long and hard about what should be obtained from a filter (like edges).
- In Deep Learning these filters are learned during the training and automatically generate new features:

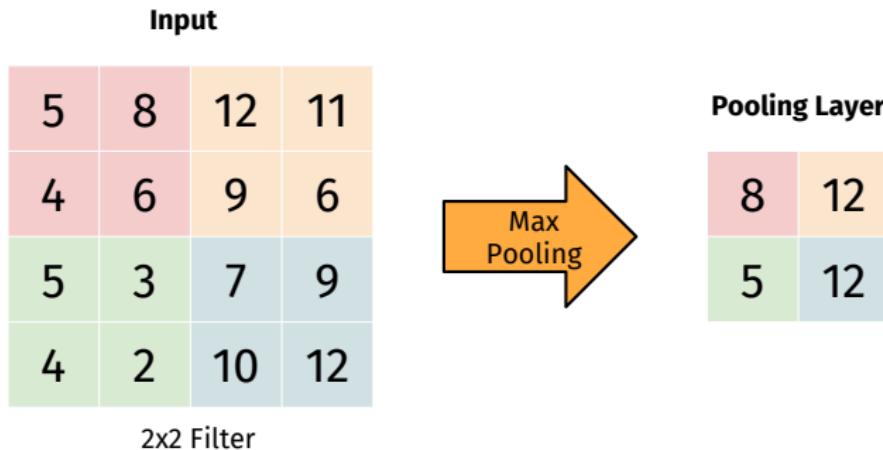


Source: Lee, H., Grosse, R., Ranganath, R., & Ng, A. Y. (2011). Unsupervised learning of hierarchical representations with convolutional deep belief networks. Communications of the ACM, 54(10), 95-103. *"Each feature can be thought of as a filter, which filters the input image for that feature (a nose). If the feature is found, the responsible unit or units generate large activations, which can be picked up by the later classifier stages as a good indicator that the class is present."*

Pooling

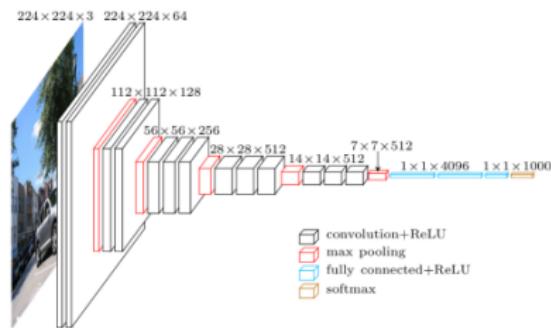
Down-sampling of images:

- Less weights to fit
 - Smaller net size (less memory usage, faster fitting process)
 - Reduces overfitting



Lets Get Deep

- The power of Deep Neural Nets comes from chaining hidden layer such as convolution layers, pooling layers, and so on
- This deep structure allows the network to create powerful features and explore complex structures within the data
- VGG16 architecture:



Source: <https://www.cs.toronto.edu/~frossard/post/vgg16/>

Pre Trained Models

| Model | Size | Parameters | Depth |
|-------------------|--------|-------------|-------|
| Xception | 88 MB | 22,910,480 | 126 |
| VGG16 | 528 MB | 138,357,544 | 23 |
| VGG19 | 549 MB | 143,667,240 | 26 |
| ResNet50 | 99 MB | 25,636,712 | 168 |
| InceptionV3 | 92 MB | 23,851,784 | 159 |
| InceptionResNetV2 | 215 MB | 55,873,736 | 572 |
| MobileNet | 16 MB | 4,253,864 | 88 |
| MobileNetV2 | 14 MB | 3,538,984 | 88 |
| DenseNet121 | 33 MB | 8,062,504 | 121 |
| DenseNet169 | 57 MB | 14,307,880 | 169 |
| DenseNet201 | 80 MB | 20,242,984 | 201 |
| NASNetMobile | 23 MB | 5,326,716 | - |
| NASNetLarge | 343 MB | 88,949,818 | - |

Source: [Keras Documentation](#)

Challenges in Deep Learning

Challenges in Deep Learning

- **Finding a Good Architecture**

Having that many possibilities of combining hidden layer, optimizer, and activation functions we run into the problem of finding a good architecture.

- Transfer learning, use already trained models, adjust them to your data situation, and train (a subset of) the weights.
- Architecture search

- **Expensive Training**

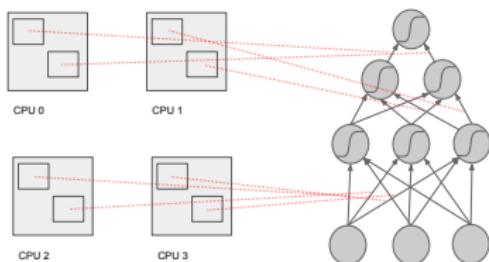
Training of DNNs require billions of simple operations, hence training **one** DNN might take weeks.

- Use GPU server for serious applications. Why? See next slides.

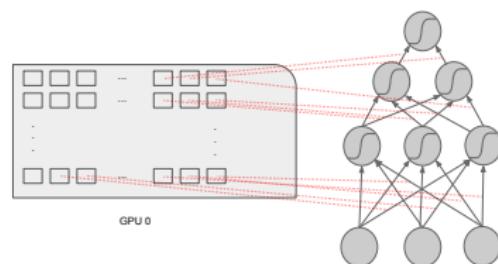
Tech Stack

Hardware

- Deep Neural Networks require special hardware to be trained efficiently.
- The training is done using Graphics Processing Units and a special programming language called **CUDA**.
- Training on standard CPUs takes a very long time and gets infeasible for anything but toy examples.



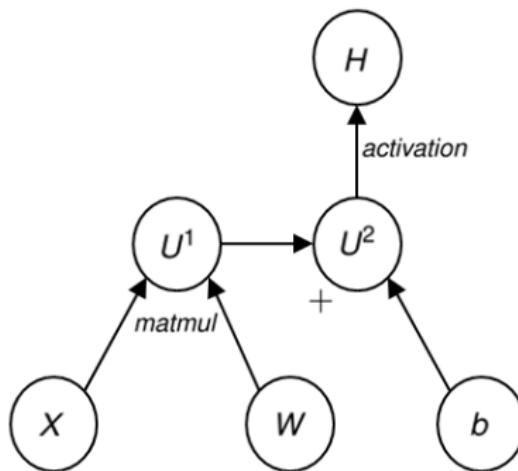
Each CPU can do 2 – 8 parallel computations.



A single GPU can do thousands of simple parallel computations.

CUDA is a very **low level** programming language and thus writing code for Deep Learning requires a lot of work. Software projects, like TensorFlow and abstract CUDA provide additional functionality.

The basic concept of calculations in deep neural networks is a *computational graph*, which describes the dependency structure of the network.



Computational graph for $f(XW + b)$.

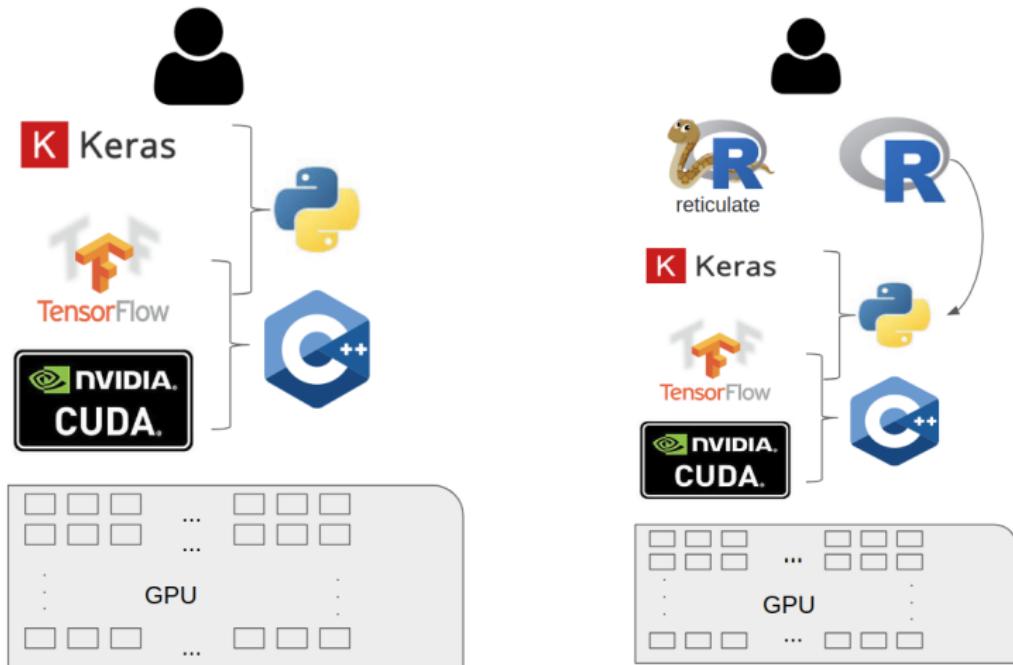


- Open-source framework developed by google.
- Rather low-level and aimed to directly work with computational graphs.
- Mainly support for Python (R support only via *reticulate*).
- Widely used and well documented.



- Open-source high-level API for Deep Learning.
- Can run on top of **TensorFlow**, CNTK or Theano.
- Mainly support for Python (R support only via **reticulate**).
- Widely used and well documented.

Keras in R



- Deep learning in R is the same as in Python.
- Communication from R to Python via **reticulate**.

- Syntax is (almost) identical.
- Same functionality.
- Same speed (slight overhead for communication between R and Python).
- More difficult to debug.

Useful resources:

<https://keras.rstudio.com/>

<https://rstudio.github.io/reticulate/articles/introduction.html>



- Open-source framework developed by facebook.
- Reimplementation of Torch.
- Only support for Python.
- Widely used and well documented.

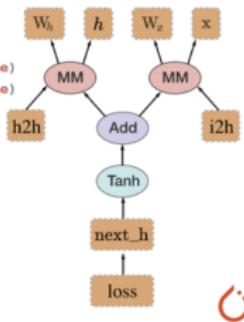
- PyTorch uses dynamic computational graphs. Instead of define and run the graph, the graph is defined by run:

A graph is created on the fly

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)

h2h = torch.mm(W_h, prev_h.t())
i2h = torch.mm(W_x, x.t())
next_h = h2h + i2h
next_h = next_h.tanh()

loss = next_h.sum()
```



Source: <https://github.com/pytorch/pytorch>

Where to Start in the DL Jungle?

Getting Started with Keras - Installation

- Install keras using pip (or pip3 for python3) from the command line:

```
# python
pip install --upgrade tensorflow
pip install keras
```

```
# python3
pip3 install --upgrade tensorflow
pip3 install keras
```

- On linux you may need to run the commands as sudo.
- The GPU version requires CUDA and cuDNN. Installation is then done by

```
pip install --upgrade tensorflow-gpu
```

Getting Started with Keras - Overview

- Instead of introducing theory, we want to get into the topic by applying it.
- We are using examples similar to the book **Deep Learning with Python** that are prepared as **notebooks**.
- **But:** When using something new, e.g. a convolution layer or optimizer, try to understand what it does and why it might be beneficial!
- First of all, we have to load keras and import models and layers:
`python import keras # equal to R's
library command from keras import models
from keras import layers`

Gettins Started with Keras - Example Data 1

```
import numpy as np
from keras.datasets import cifar10 # Our dataset
# See https://www.cs.toronto.edu/~kriz/cifar.html

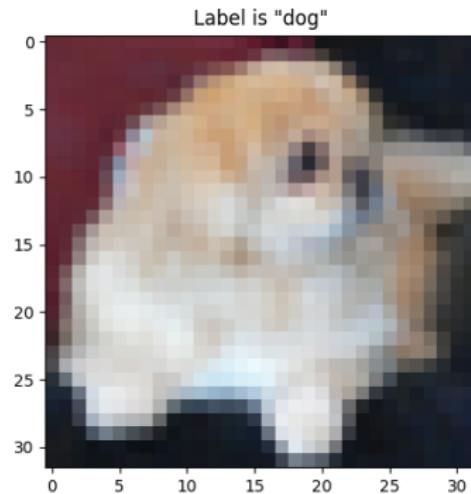
(X_train, Y_train), (X_test, Y_test) = cifar10.load_data()
X_train.shape
## (50000, 32, 32, 3)

categories = ["airplane", "automobile", "bird", "cat",
              "deer", "dog", "frog", "horse", "ship",
              "truck"]

import matplotlib.pyplot as plt

category = categories[np.asscalar(Y_train[2305])]
plt.title('Label is \"{}\"'.format(label=category))
plt.imshow(X_train[2305], cmap='gray')
plt.show()
```

Getting Started with Keras - Example Data 2



Getting Started with Keras - Example Data 3

The images are represented as rgb colors, each of the 32×32 pixels is represented as three numbers between 0 and 255. For the training we want to use values between 0 and 1:

```
train_images = X_train.astype('float32') / 255
test_images = X_test.astype('float32') / 255
```

Additionally, we have to convert the labels to a categorical data type:

```
from keras.utils import to_categorical

train_labels = to_categorical(Y_train)
test_labels = to_categorical(Y_test)
```

Getting Started with Keras - DNN 1

```
network = models.Sequential()

network.add(layers.Conv2D(64, (3, 3),
                       input_shape=(32, 32, 3), padding='same'))
# now: model.output_shape == (None, 64, 32, 32)

network.add(layers.Flatten())

# Add fully connected hidden layer:
network.add(layers.Dense(units=128,
                        activation='relu'))

# Add output layer which maps each category to a neuron:
network.add(layers.Dense(10, activation='softmax'))

# Make the network ready for training:
network.compile(optimizer='rmsprop',
                 loss='categorical_crossentropy',
                 metrics=['accuracy'])
```

Getting Started with Keras - DNN 2

```
network.summary()
## -----
## Layer (type)                  Output Shape             Param #
## =====
## conv2d_2 (Conv2D)            (None, 32, 32, 64)      1792
## -----
## flatten_2 (Flatten)          (None, 65536)           0
## -----
## dense_3 (Dense)              (None, 128)            8388736
## -----
## dense_4 (Dense)              (None, 10)              1290
## -----
## Total params: 8,391,818
## Trainable params: 8,391,818
## Non-trainable params: 0
## -----
```

Getting Started with Keras - DNN 3

```
# Train network:  
network.fit(train_images, train_labels, epochs=20, batch_size=128)  
  
network.evaluate(test_images, test_labels)  
## [3.585078500747681, 0.4522]
```

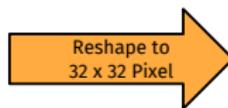
Getting Started with Keras - DNN 4

Now we try to predict the class of a new image. Therefore, we have to rescale the original image to the same input shape of the training images:

Original Image of Wuschel
(a cat)

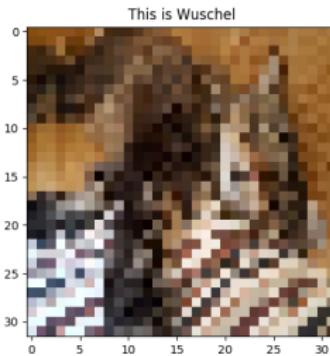


Rescaled Image of Wuschel
(still a cat, but very blurred)



Getting Started with Keras - DNN 5

```
from keras.preprocessing.image import load_img, img_to_array  
img = load_img('wuschel.jpg', target_size=(32, 32))  
img = img_to_array(img).astype('float32') / 255  
  
plt.title('This is Wuschel')  
plt.imshow(img, cmap='gray')  
plt.show()
```



Getting Started with Keras - DNN 6

```
wuschel_pred = network.predict(np.array([img]))
wuschel_pred
## array([[0.01510279, 0.03711035, 0.04978644, 0.25755176,
##         0.03696484, 0.23572417, 0.00041433, 0.3290789 ,
##         0.00335977, 0.03490658]], dtype=float32)

categories[wuschel_pred.argmax()]
## 'horse'
```

Getting Started with Keras - Transfer Learning 1

```
from keras.applications import VGG16

conv_base = VGG16(weights='imagenet',
                  include_top=False,
                  input_shape=(32, 32, 3))

from keras import optimizers

model = models.Sequential()
model.add(conv_base)
model.add(layers.Flatten())
model.add(layers.Dense(256, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

conv_base.trainable = False

model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Getting Started with Keras - Transfer Learning 2

```
model.summary()
## -----
## Layer (type)                  Output Shape             Param #
## =====
## vgg16 (Model)                (None, 1, 1, 512)      14714688
## -----
## flatten_3 (Flatten)           (None, 512)            0
## -----
## dense_5 (Dense)               (None, 256)            131328
## -----
## dense_6 (Dense)               (None, 10)             2570
## -----
## Total params: 14,848,586
## Trainable params: 133,898
## Non-trainable params: 14,714,688
## -----
```

Getting Started with Keras - Transfer Learning 3

```
# Train network:  
model.fit(train_images, train_labels, epochs=20, batch_size=128)  
  
model.evaluate(test_images, test_labels)  
## [1.1461146575927734, 0.6175]
```

Getting Started with Keras - Transfer Learning 5

```
wuschel_pred = model.predict(np.array([img]))
wuschel_pred
## array([[1.4046730e-03, 1.0323318e-03, 1.2766185e-01,
##         4.6453097e-01, 5.1652599e-02, 5.0368346e-02,
##         2.9299492e-01, 1.0088098e-02, 2.4648944e-05,
##         2.4157017e-04]], dtype=float32)

categories[wuschel_pred.argmax()]
## 'cat'
```

What Next?

Improving Performance:

- Choosing the number of epochs using early stopping
- Data augmentation (increasing size of training data by rotating, scaling, flipping, . . .)

More Difficult Tasks:

- RNN, LSTM, GAN, Reinforcement Learning
- Natural Language Processing